

# Linear Regression

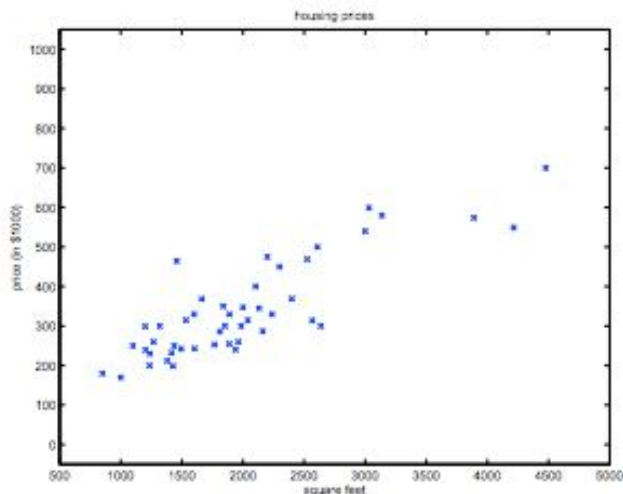
---

GA DAT3

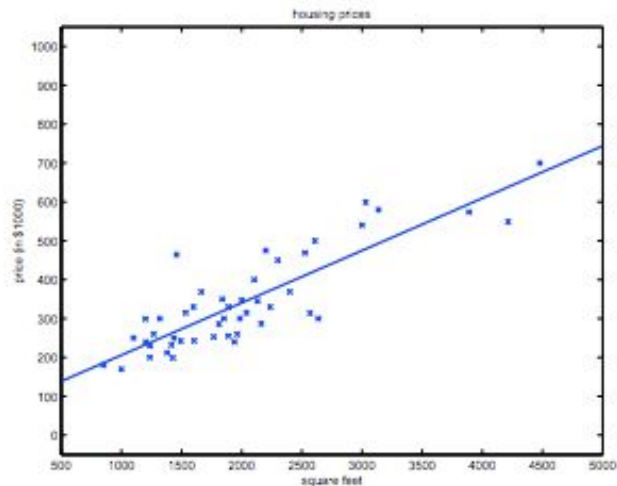
# Agenda

1. Overview
2. Cost Function
3. Gradient Descent
4. Normal Equation
5. Probabilistic Interpretation
6. Locally Weighted Linear Regression

# Linear Regression Overview



Supervised Learning



Regression: predict real-value  
output

Classification: predict discrete value  
output


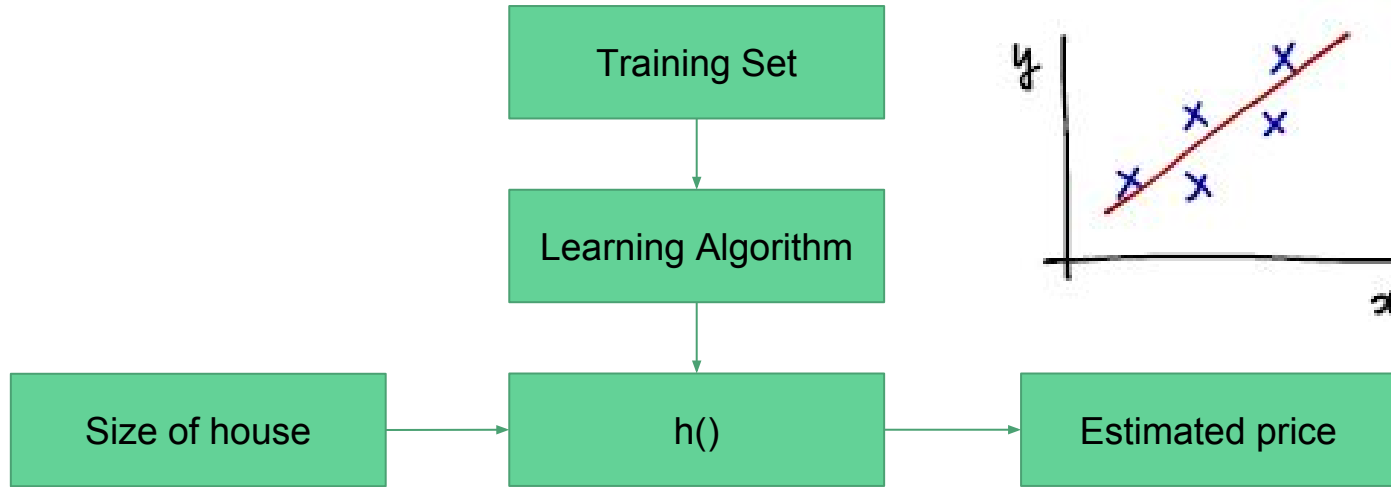
Training set of housing prices (Portland, OR)	Size in feet <sup>2</sup> ( $x$ )	Price (\$) in 1000's ( $y$ )
	2104	460
	1416	232
	1534	315
	852	178
	...	...

Notation:

$m$  = Number of training examples

$x$ 's = "input" variable / features

$y$ 's = "output" variable / "target" variable




Linear regression with one variable :  
**Univariate linear regression**

# Cost Function : Squared Error Loss

Minimize: 
$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2.$$

WHY?



# General Multivariate Form

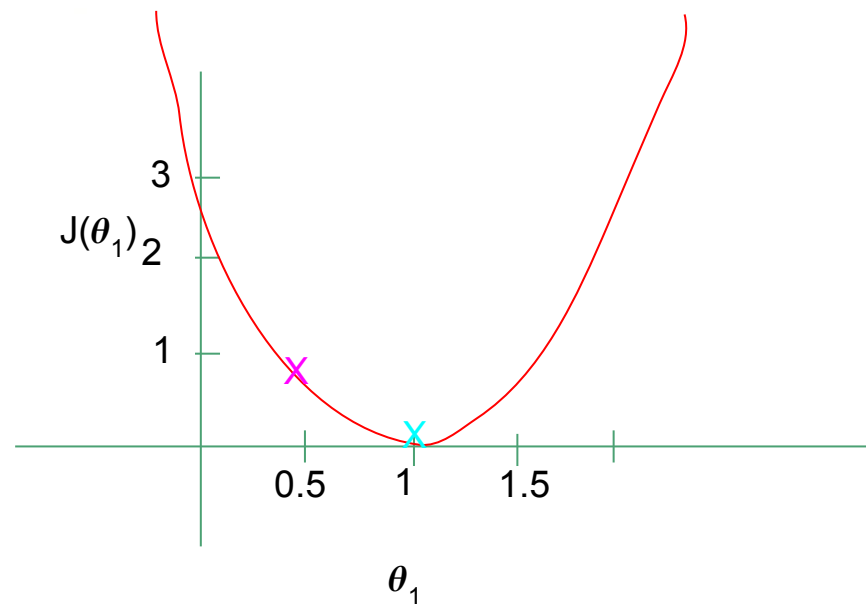
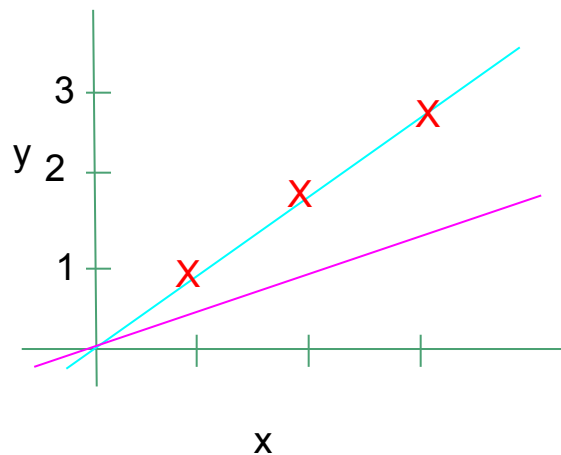
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x$$

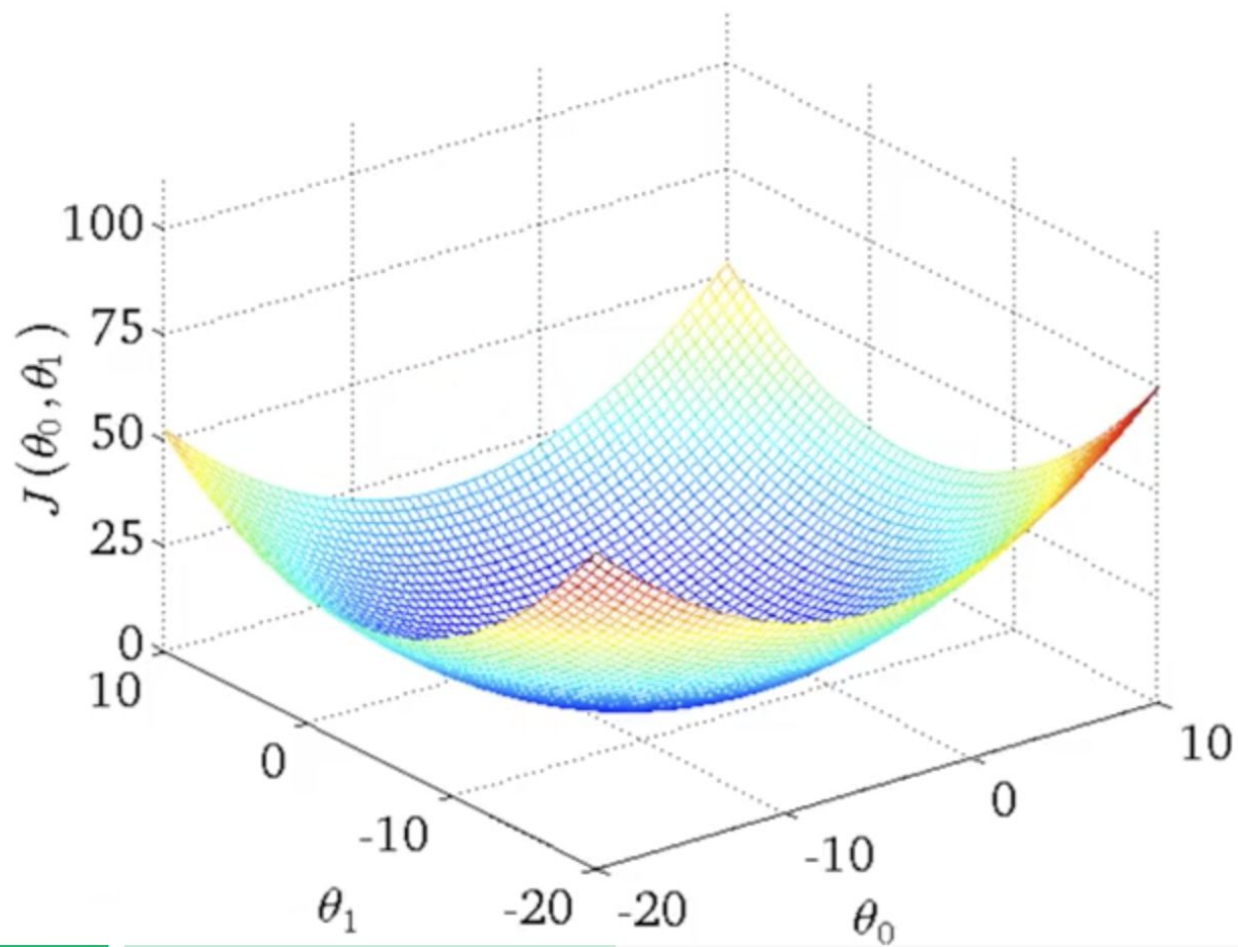
# Cost Function : Squared Error Loss

Minimize:

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2.$$







# Gradient Descent

## Update Rule:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta).$$

Learning rate

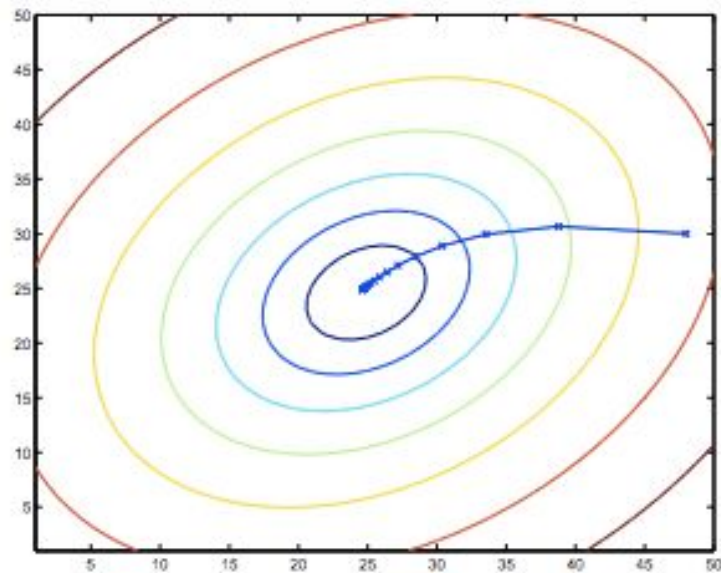
Repeat until convergence {

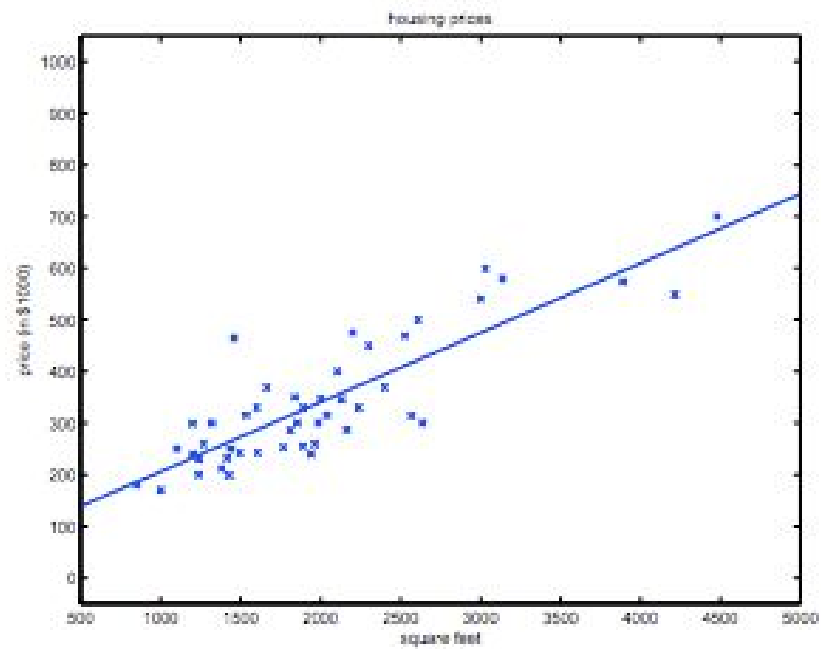
$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

} For sufficiently small  $\alpha$ ,  $\theta$  should decrease on every iteration.  
But if  $\alpha$  is too small, gradient descent can be slow to converge."

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\ &= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y) \\ &= (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} \left( \sum_{i=0}^n \theta_i x_i - y \right) \\ &= (h_{\theta}(x) - y) x_j \end{aligned}$$

What are the axes?





# Batch vs Stochastic Descent

for  $\theta$ :

for each example:

...

for each example:

for  $\theta$ :

...

# Normal Equations

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \frac{1}{2} (X\theta - \vec{y})^T (X\theta - \vec{y}) \\&= \frac{1}{2} \nabla_{\theta} (\theta^T X^T X \theta - \theta^T X^T \vec{y} - \vec{y}^T X \theta + \vec{y}^T \vec{y}) \\&= \frac{1}{2} \nabla_{\theta} \text{tr} (\theta^T X^T X \theta - \theta^T X^T \vec{y} - \vec{y}^T X \theta + \vec{y}^T \vec{y}) \\&= \frac{1}{2} \nabla_{\theta} (\text{tr} \theta^T X^T X \theta - 2 \text{tr} \vec{y}^T X \theta) \\&= \frac{1}{2} (X^T X \theta + X^T X \theta - 2 X^T \vec{y}) \\&= X^T X \theta - X^T \vec{y}\end{aligned}$$

$$\theta = (X^T X)^{-1} X^T \vec{y}.$$

# Gradient Descent vs Normal Equation

## Gradient Descent

Need to choose  $\alpha$

Need many iterations

Work well even with large  $n$

## Normal Equation

No need to choose  $\alpha$

Don't need to iterate

Need to compute  $(X^T X)^{-1}$

Slow if  $n$  is large

# Probabilistic Interpretation

$$y^{(i)} = \theta^T x^{(i)} + \epsilon^{(i)}$$

$$p(\epsilon^{(i)}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(\epsilon^{(i)})^2}{2\sigma^2}\right).$$


We can minimise **cost**  
or  
**maximise likelihood**

What's the likelihood?



# Likelihood

This is a **conditional probability density**


$$\begin{aligned} L(\theta) &= \prod_{i=1}^m p(y^{(i)} \mid x^{(i)}; \theta) \\ &= \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \end{aligned}$$

Can you write this as a sum?


## Likelihood - For you advanced folks

$$\begin{aligned}\ell(\theta) &= \log L(\theta) \\ &= \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\ &= \sum_{i=1}^m \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\ &= m \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \cdot \frac{1}{2} \sum_{i=1}^m (y^{(i)} - \theta^T x^{(i)})^2.\end{aligned}$$

Hence, maximizing  $\ell(\theta)$  gives the same answer as minimizing

$$\frac{1}{2} \sum_{i=1}^m (y^{(i)} - \theta^T x^{(i)})^2,$$

# Locally Weighted Linear Regression



classic tradeoff

# Locally Weighted Linear Regression

Instead of doing:

1. Fit  $\theta$  to minimize  $\sum_i (y^{(i)} - \theta^T x^{(i)})^2$ .
2. Output  $\theta^T x$ .

We do:

1. Fit  $\theta$  to minimize  $\sum_i w^{(i)} (y^{(i)} - \theta^T x^{(i)})^2$ .
2. Output  $\theta^T x$ .

# Locally Weighted Linear Regression

Good examples of weights?

$$w^{(i)} = \exp \left( -\frac{(x^{(i)} - x)^2}{2\tau^2} \right)$$

# Locally Weighted Linear Regression

Parametric vs Non-Parametric

What was k-Nearest Neighbours?

[http://learning.cis.upenn.edu/cis520\\_fall2009/index.php?n=Lectures.LocalLearning#toc8](http://learning.cis.upenn.edu/cis520_fall2009/index.php?n=Lectures.LocalLearning#toc8)

Q??

---